

# HPC Concepts

...

Queues Schedulers, and workflow design for shared systems.

# Overview

- Types
  - Of Computing systems
  - Of Jobs
- Limitations
  - Compute
  - I/O
- Cluster components
- Seperation
- Queuing
  - Queue Partitioning
  - Queue Filling
- Workflow organization as solution
- Examples

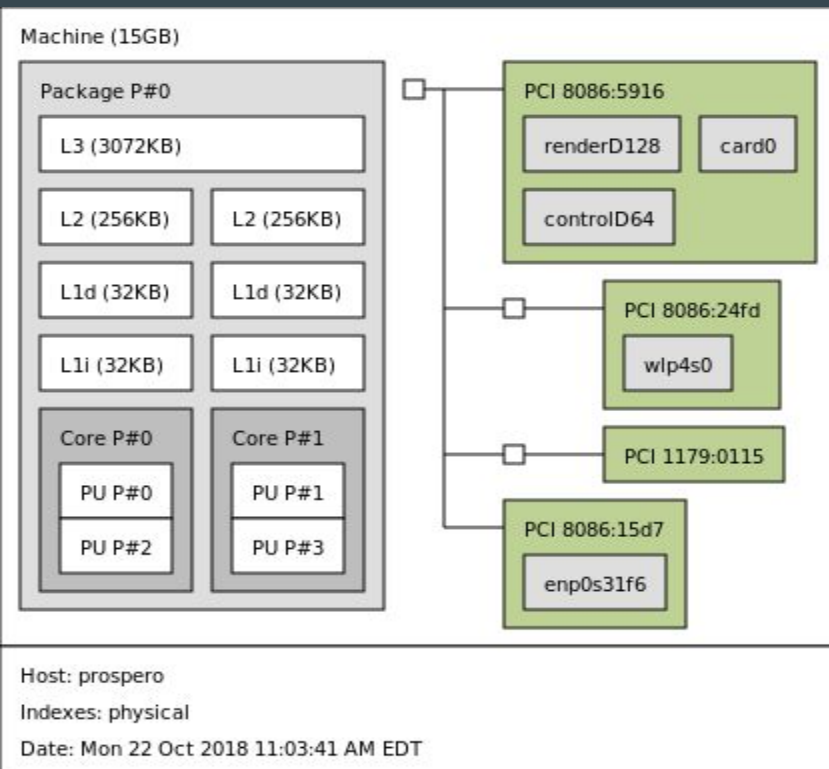
# **What is a HPC?**

# **What is a HPC?**

**a device for turning compute-bound problems into I/O-bound problems.**

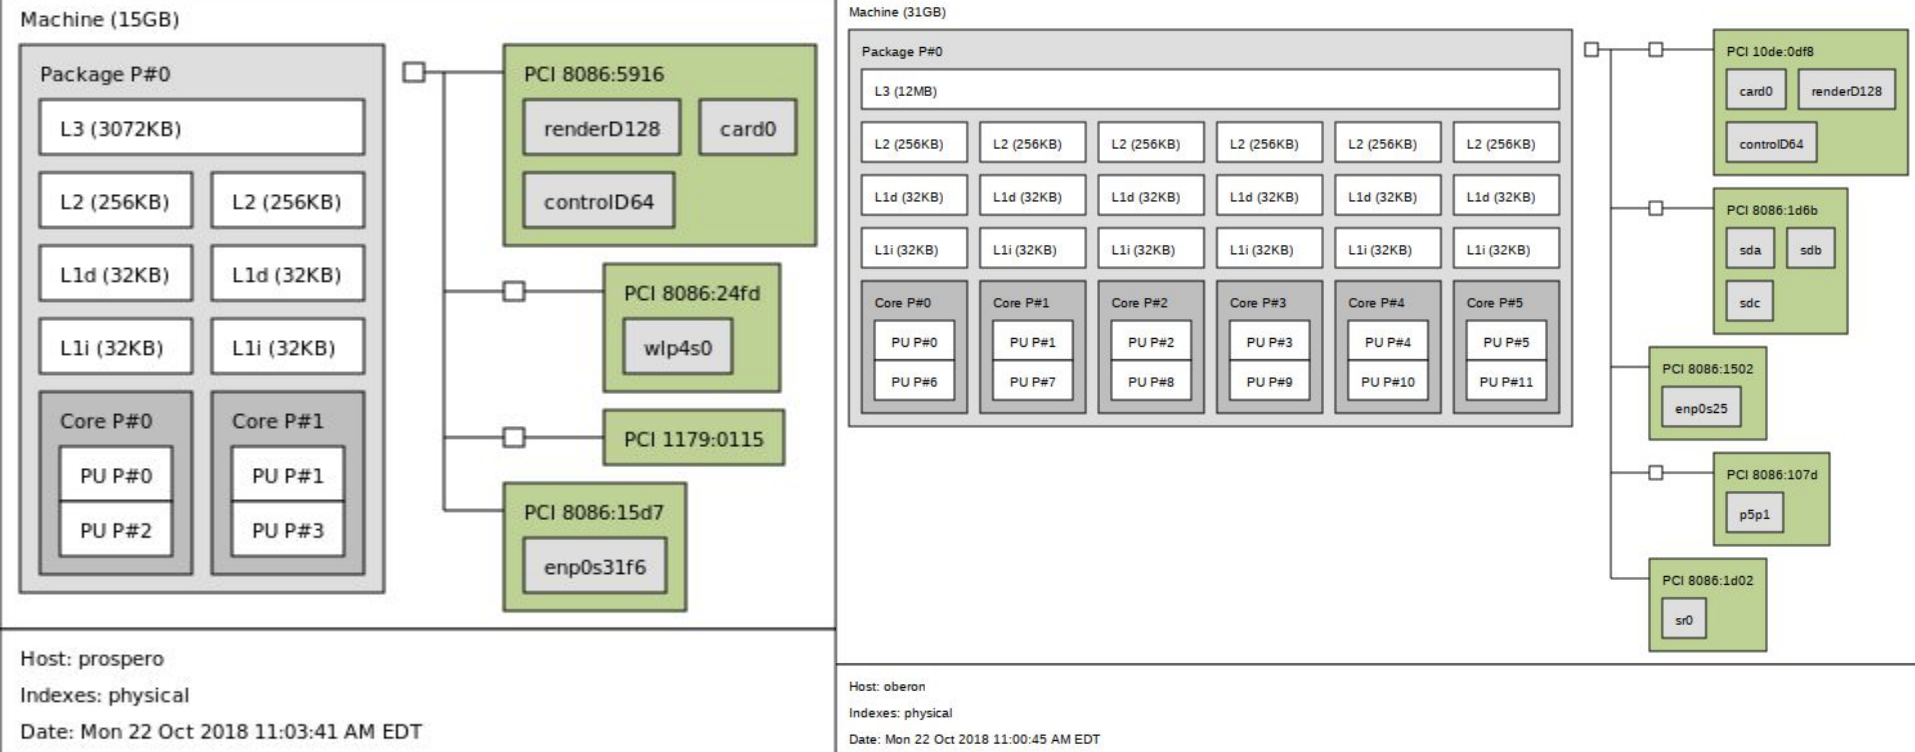
**– Ken Batcher**

# More powerful CPU's, More Chips Per board



'Business Laptop'

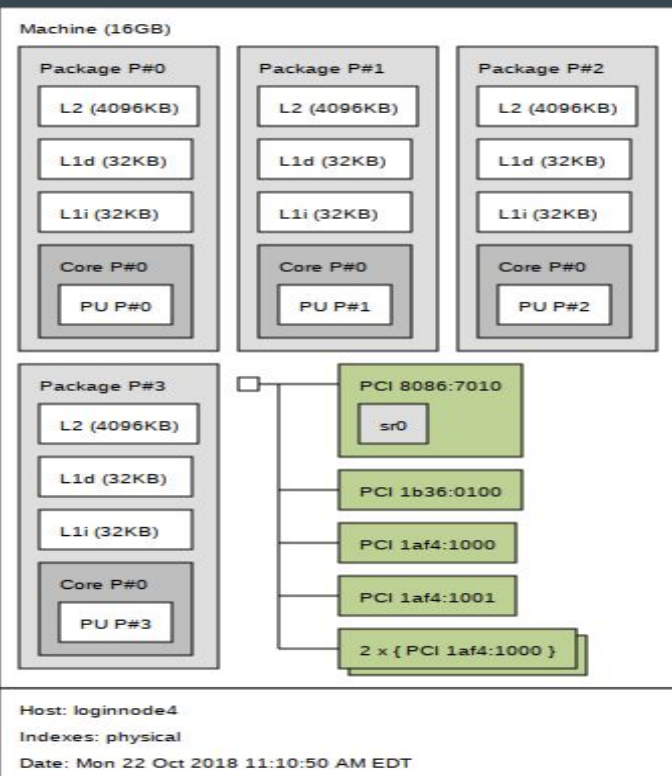
# More powerful CPU's, More Chips Per board



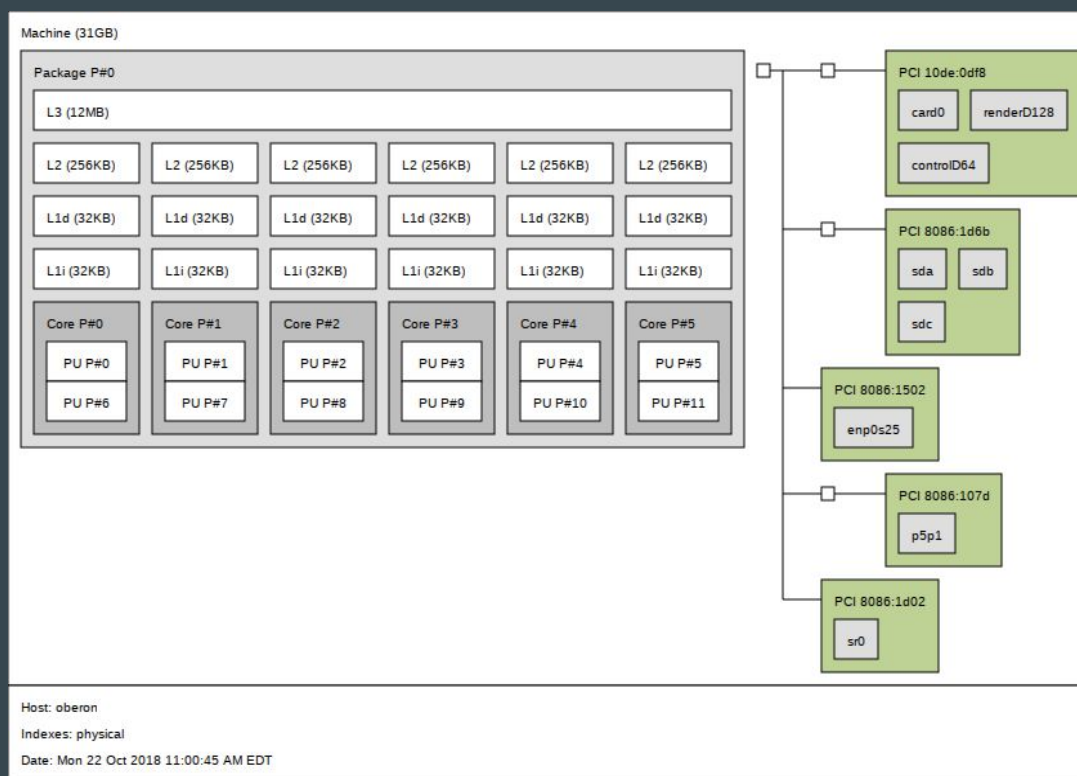
'Business Laptop'

'Mid-level Workstation'

# More powerful CPU's, More Chips Per board - Reorganized



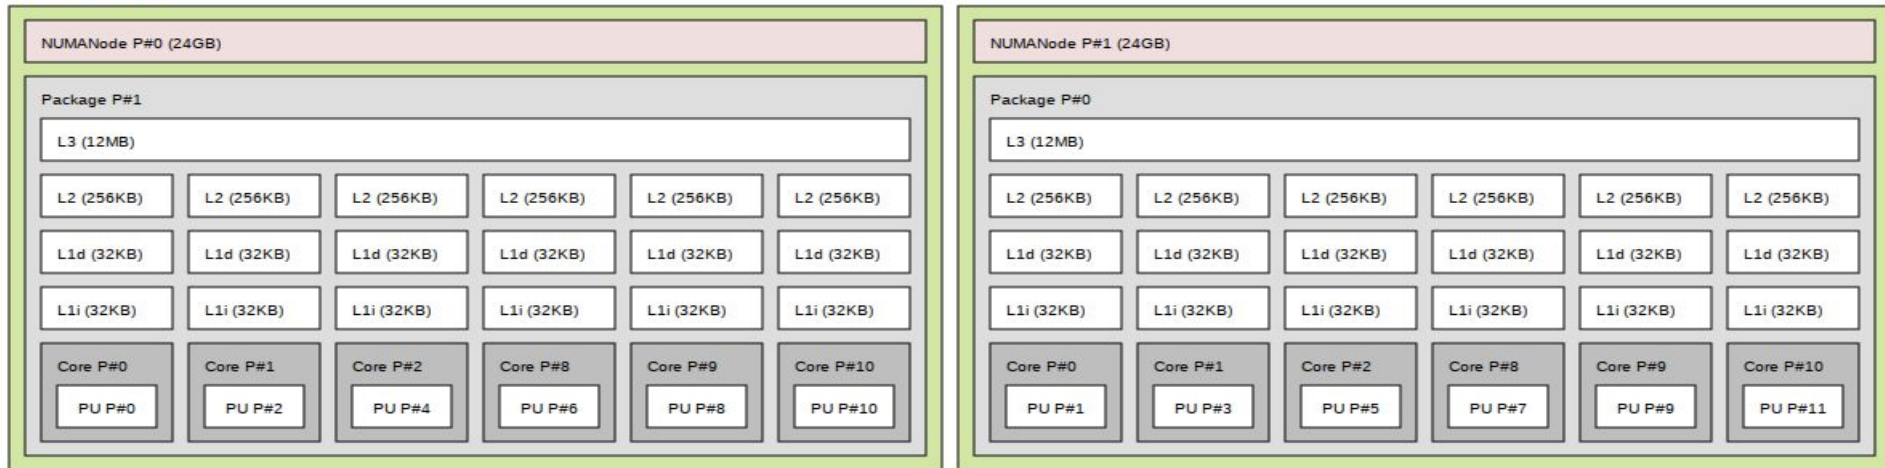
Head node



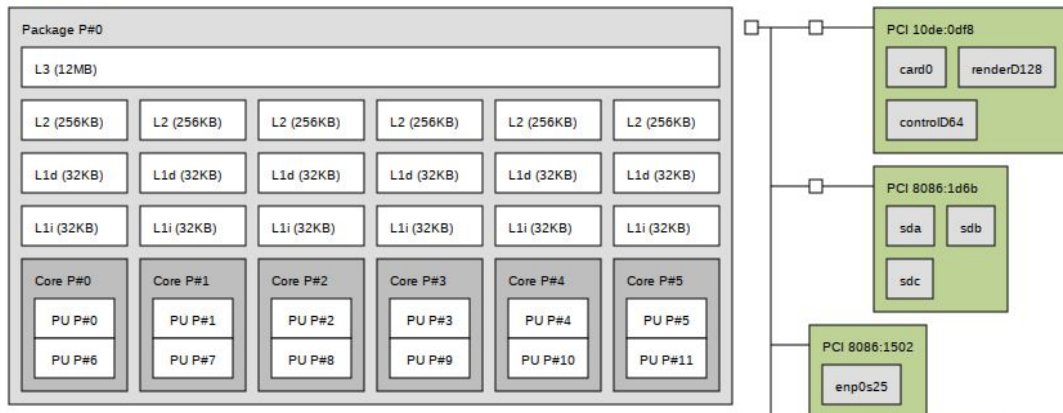
'Mid-level Workstation'

# More powerful CPU's, More Chips Per board - Reorganized

Machine (48GB total)



Machine (31GB)

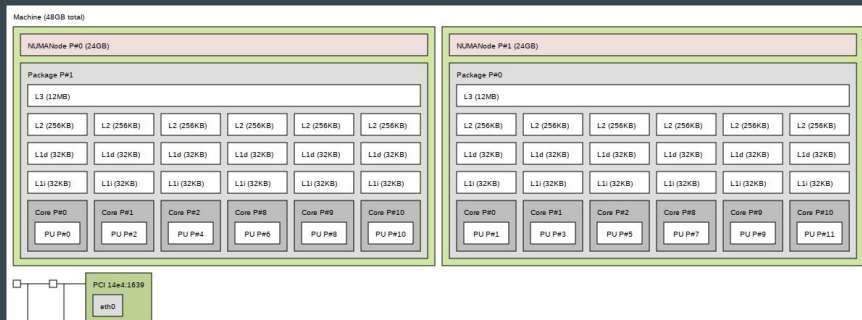


Small worker node

Same  
'Mid-level Workstation'



# More powerful CPU's, More Chips Per board - Reorganized



Small worker node

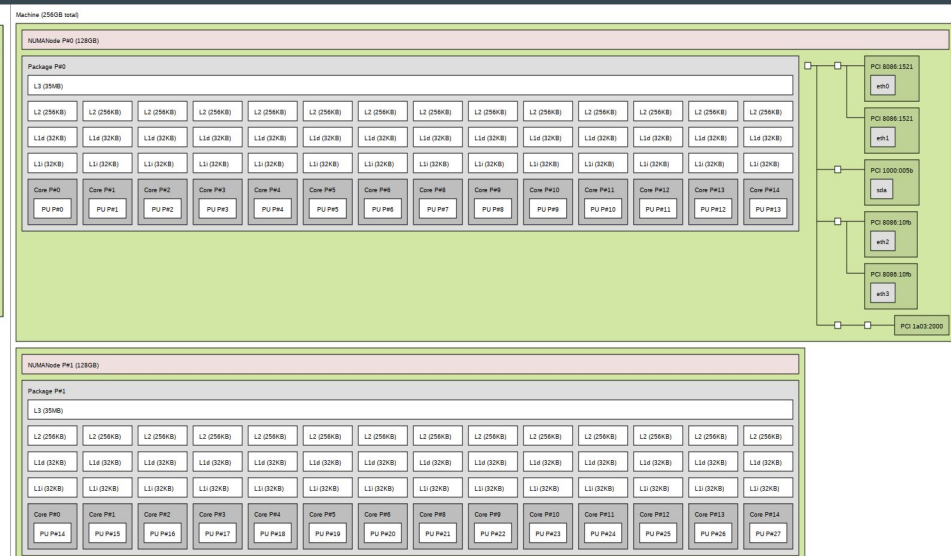
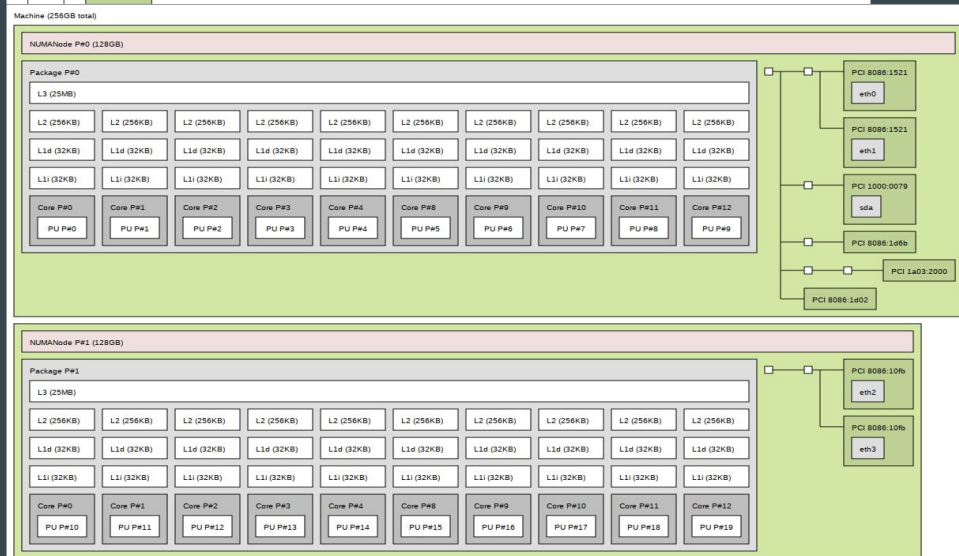
2x 6 cores  
2x 24Gb RAM

Mid-size worker node

2x 12 cores  
2x 128 Gb RAM

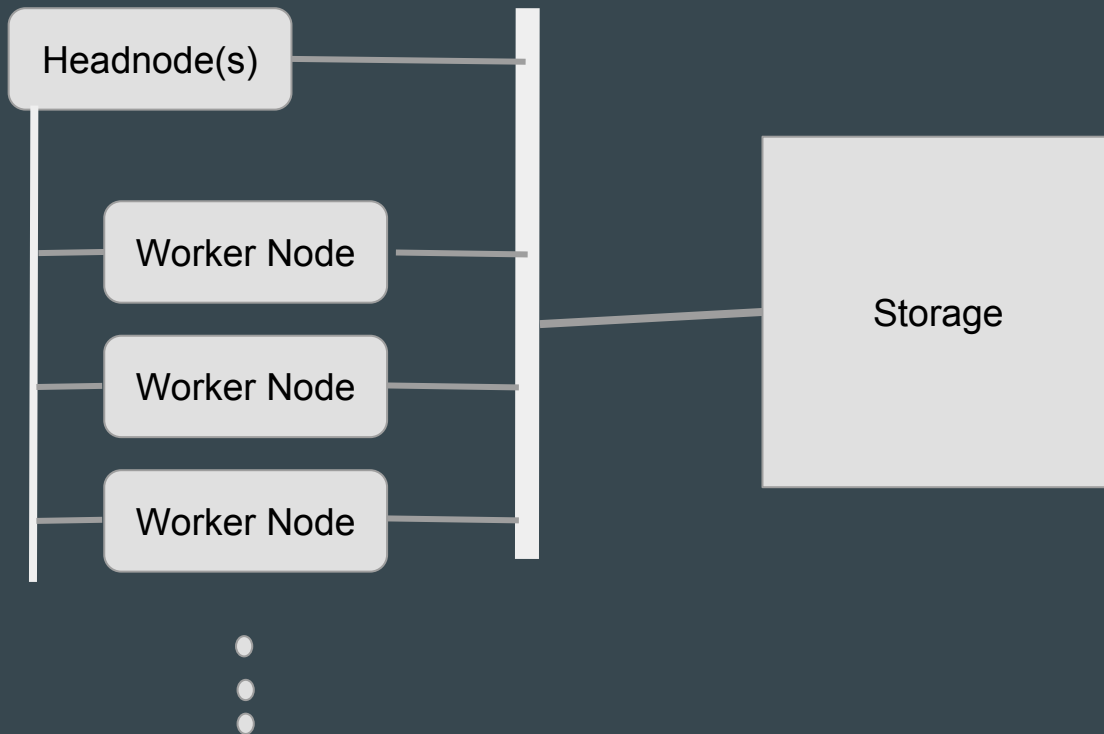
Big-compute worker node

2x 14 cores  
2x 128 Gb RAM




# More Interconnects

Private  
LAN



Types of HPC

# More Interconnects


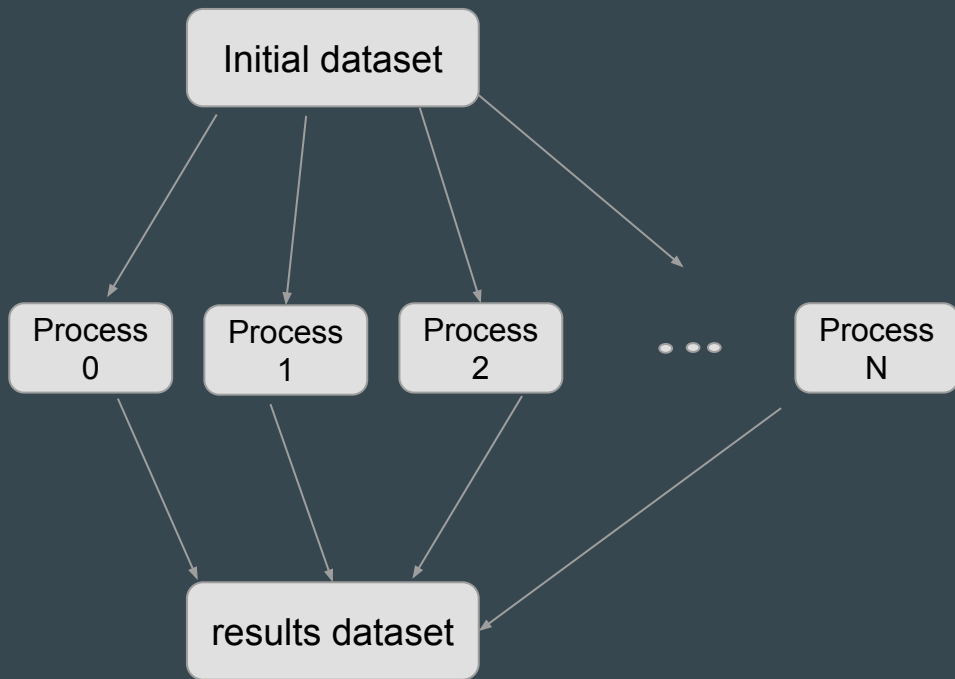


Types of HPC

# More Interconnects

“Pleasingly Parallel”


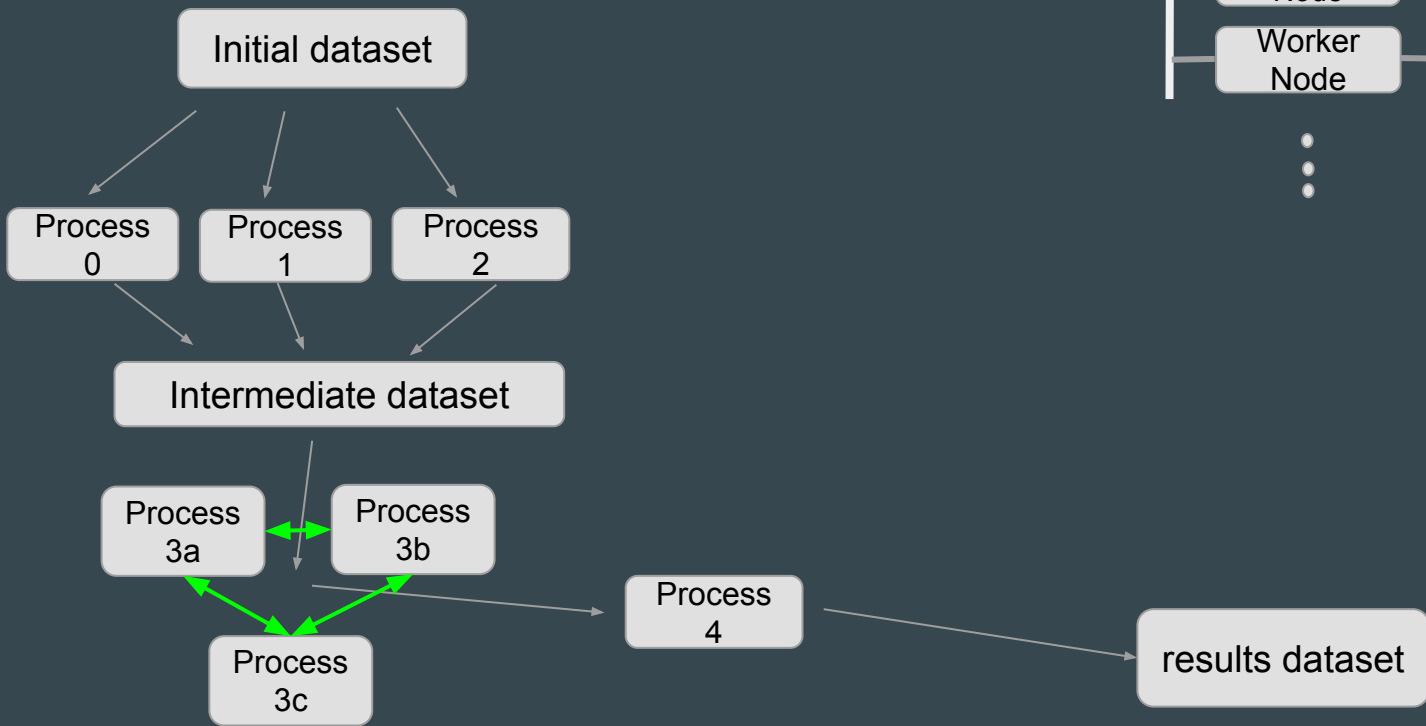
Minimal complexity



Types of HPC

# More Interconnects



Loosely coupled workflows



Types of HPC

# More Interconnects

Loosely coupled workflows



Loosely coupled jobs have dependent processes, but distribute the data according to the task they are

Dependent in terms of task order and



Independent internal to the task and its individual data flow

Special setup can improve speed but is not strictly required

Types of HPC

# More Interconnects

Tightly coupled workflows





Tightly coupled jobs have interdependent processes, and Shared memory.

Types of HPC

# More Interconnects

Tightly coupled workflows



Tightly coupled jobs have interdependent processes, and Shared memory.

Requires special physical set up and programming (PVM or MPI)



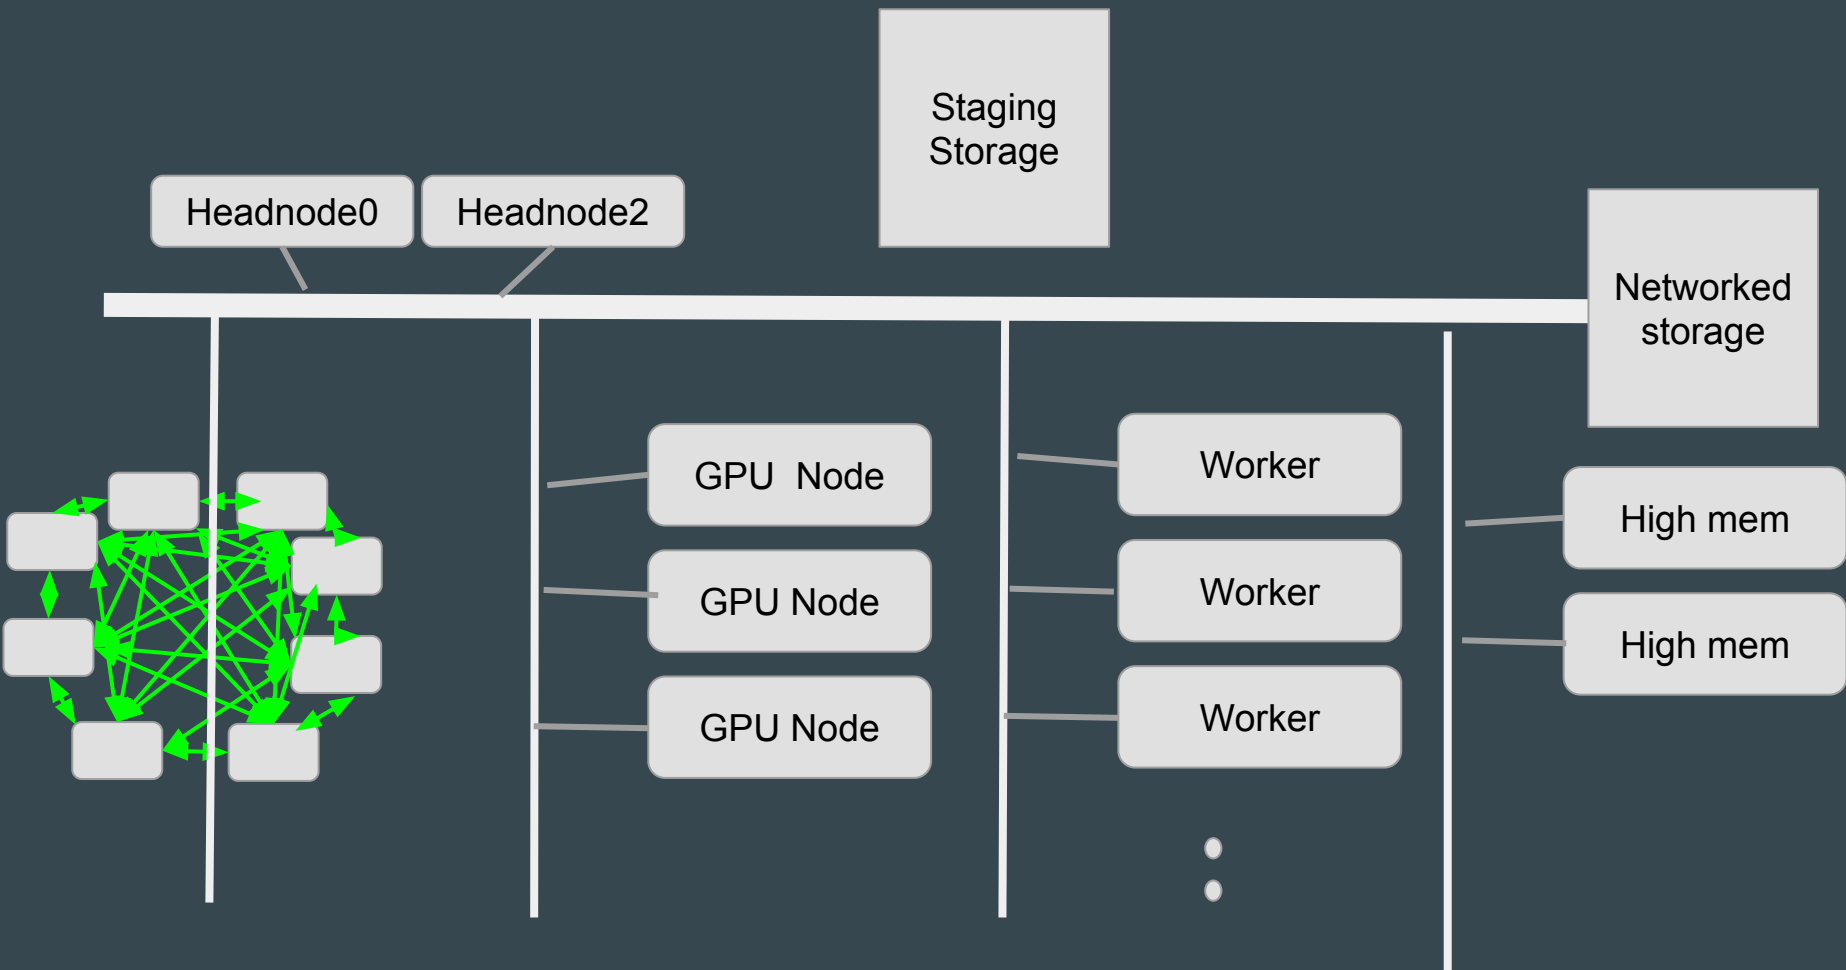
Types of HPC


# Acceleration

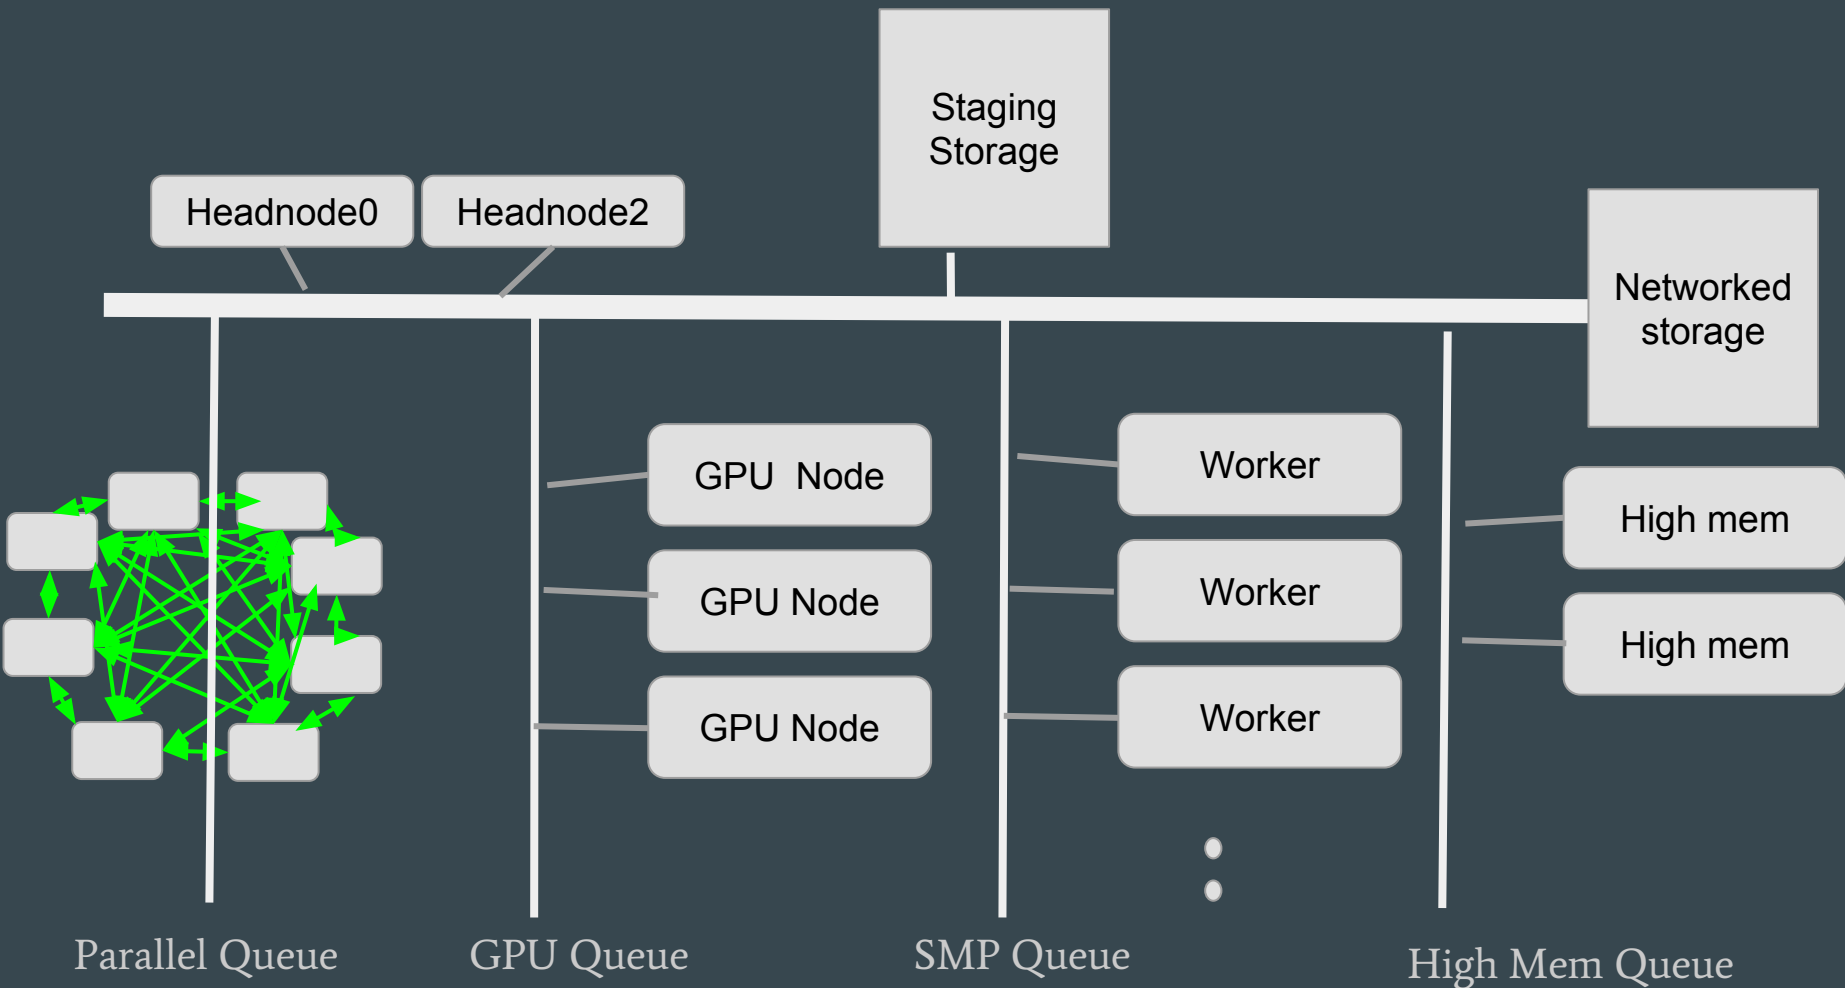
GPU's: essentially a dense tightly bound system.


ARM: Single Instruction Multiple Data (SIMD) vector processors

FPGA: Field Programmable Gate arrays  
Customizing the physical structure of the chip for specific problem sets.









# A mile high View

- Solutions are Organizational
- Matches the hardware with
  - Data flow
  - Computational patterns.

---

# Common Queue Scheduler systems

## For distributed work

- HTCondor
- Yarn

## For localized systems

- Grid Engine (SGE, UGE)
- PBS/torque
- Slurm

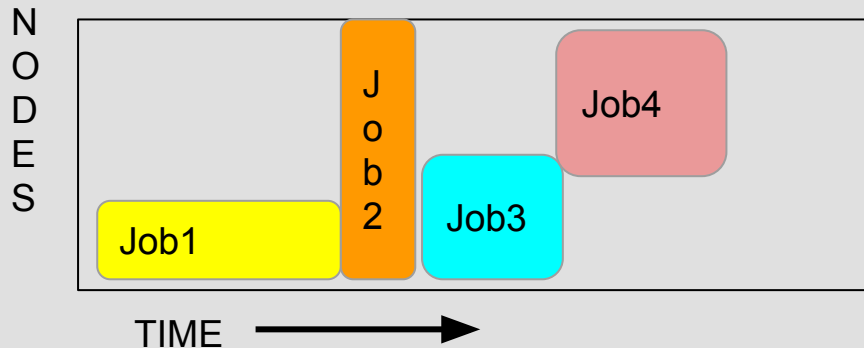
# Queue handling

- First in First Out FIFO
- Priorities (Backfill scheduling)
- Preemptions



# Queue handling

First in First Out    FIFO




# Queue handling

First in First Out    FIFO


BackFill

N  
O  
D  
E  
S



TIME →


N  
O  
D  
E  
S




TIME →

# Queue handling

## BackFill



## Optimizing backfill



# Fair share

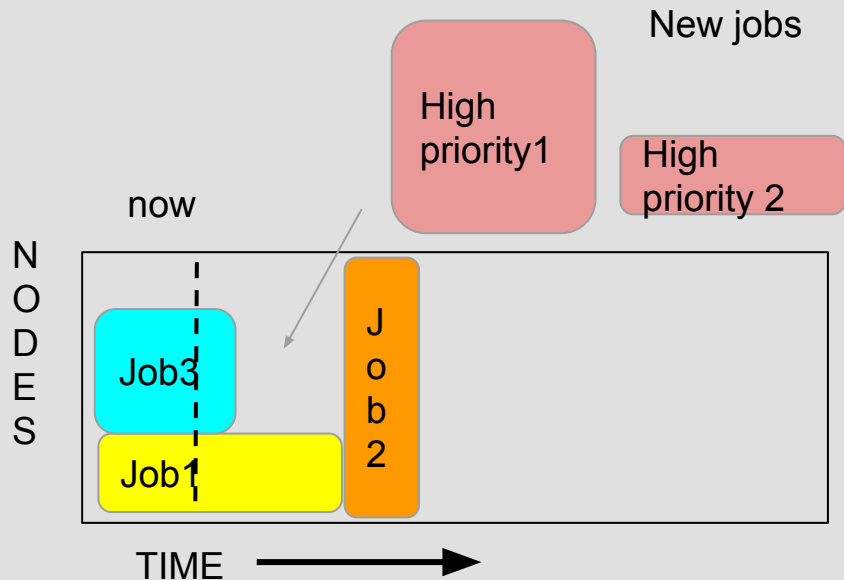
## Adds Priorities to backfill scheduling

Weight several factors to generate a priority level for each job


- How long it has been waiting
- Job Size
- Where the job is being run
- User's base priority level

# Queue handling

## Preemption



Jobs 1 & 3 are stopped and restarted



# What is Checkpointing?

- Saves a representation of the current state at regular intervals
- Allows Automation
  - Scheduler with automatically restarts
  - Add a function at the beginning to check for the save state file
    - If a save state file exists reads it in and starts from there.

```
If statefile exists:  
    Startjobat(statefile)  
else:  
    Startjobat(beginning)
```

# When to use Checkpointing

- Long running jobs
  - To not start over if there's a “hiccup”
    - Balancing between speed and loss
      - Writing to a file is much much slower than anything else.
      - Longer period between checkpoint means faster runs
      - Potential lost time with a restart is bigger.
  - Runtime is longer than maximum job time

# Planning workflows

- Seperate jobs
- Identify dependencies
- Decided what (if any) intermediate data to keep
- Match your pattern to available resources and limitations.

---



Turning compute-bound problems into I/O-bound problems.


**SMP and NUMA**

Turning compute-bound problems into I/O-bound problems.

## SMP and NUMA

---

### Shared Memory Processing *Uniform access*




Turning compute-bound problems into I/O-bound problems.

## SMP and NUMA

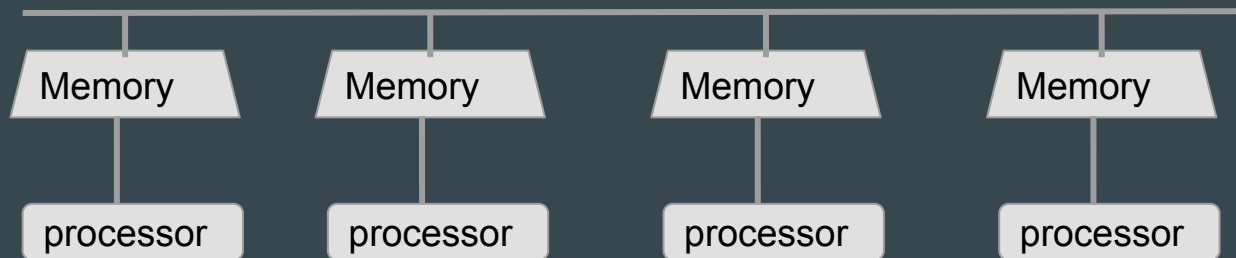
---

### Shared Memory *Uniform access*




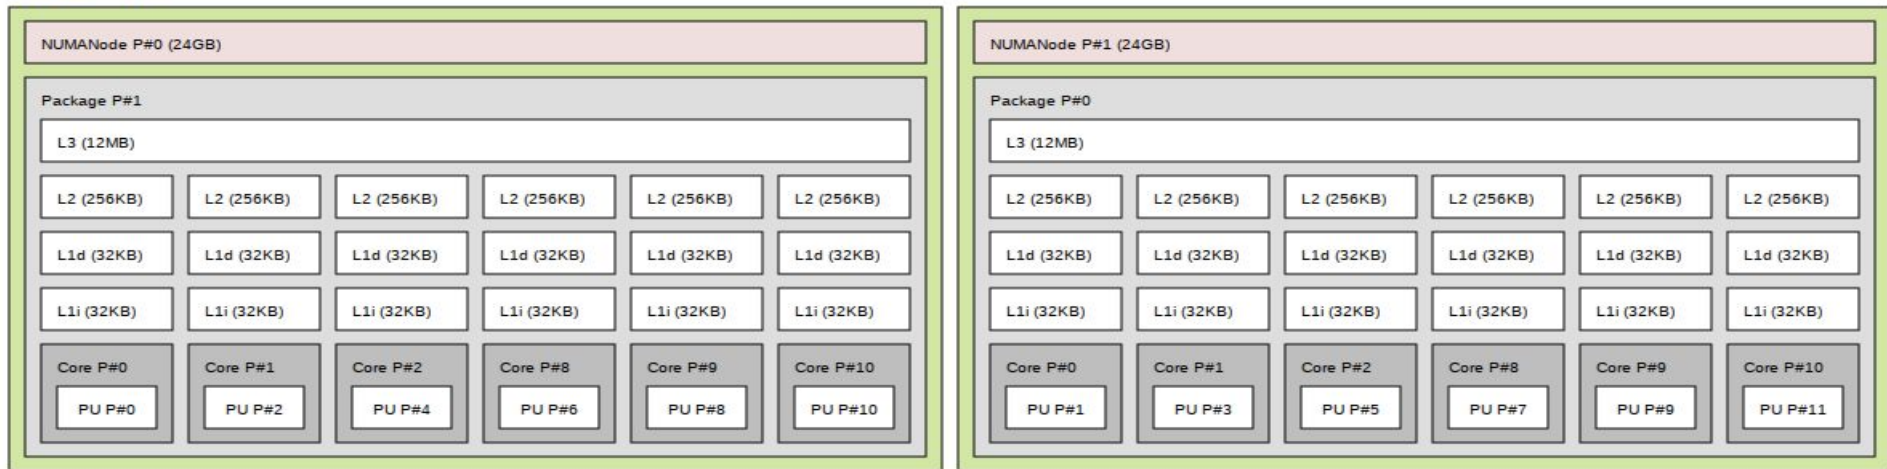
---

### Shared Memory *Non-Uniform Memory Access*

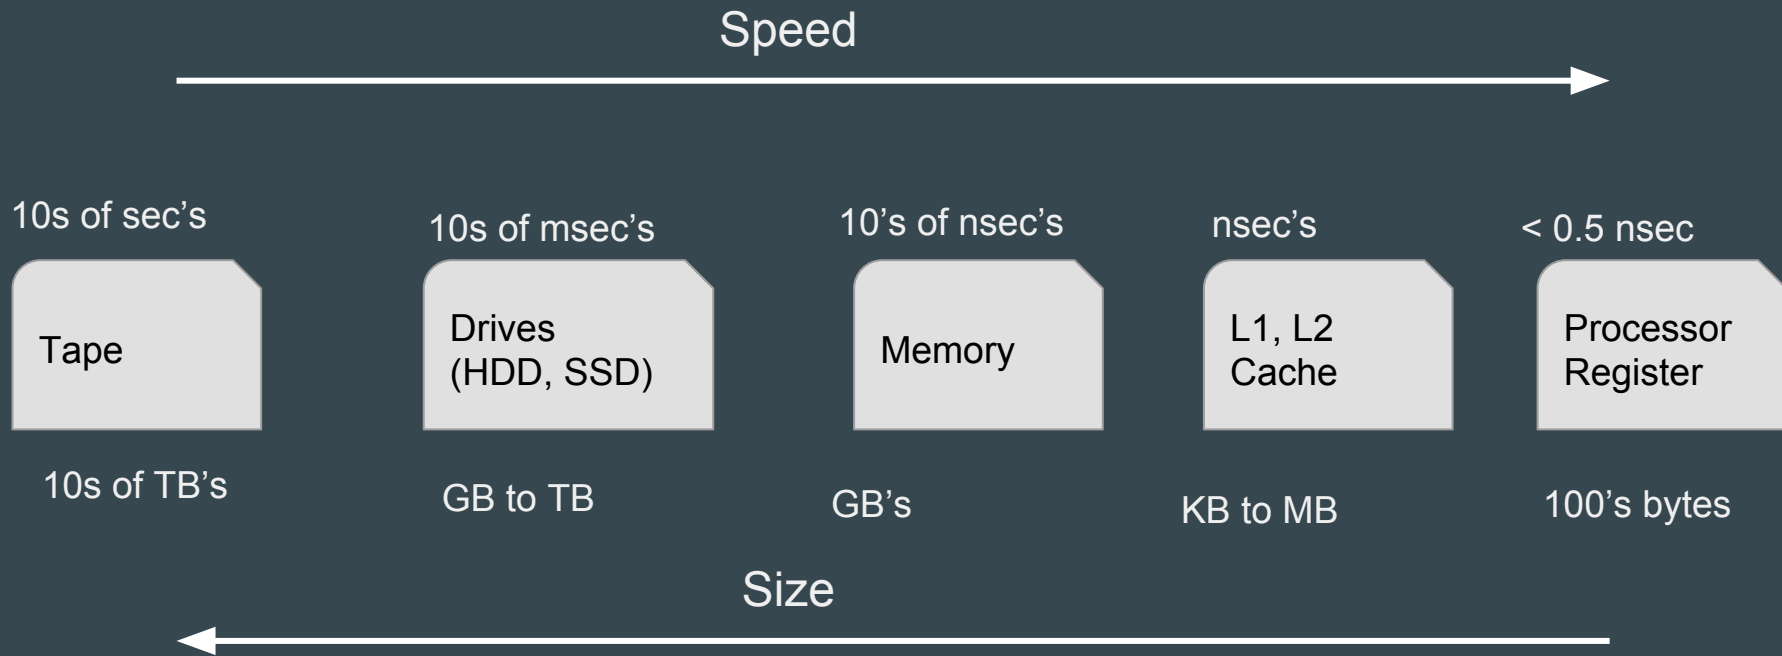


# Turning compute-bound problems into I/O-bound problems. Everything is a graph

Machine (48GB total)



# Turning compute-bound problems into I/O-bound problems.



# Distance = 1/Speed

Remote networked storage	Slow - not for computing
Nearby networked/warm storage	Pre-staging Completed data
'Fast scratch' (or /tmp if not available)	Staging, interprocess holding, bigger than RAM for for worker I/o
Tmpfs (memory as filesystem)	For faster I/o

# Distance = 1/Speed

# Caution every system is different

Remote networked storage	Slow - not for computing
Nearby networked/warm storage	Pre-staging Completed data
'Fast scratch' (or /tmp if not available)	Staging, interprocess holding, bigger than RAM for for worker I/o
Tmpfs (memory as filesystem)	For faster I/o

Check with your local systems documentation

Or ask the system folks.

# Modules

## Modules

modules allow us to keep the working space clean, when there are a lot of installed programs.

### **\$module avail**

*Will list the currently available modules*

### **\$module load <modulename>**

*Will make the module available  
Be explicit!*

### **\$module unload <modulename>**

*Will unload the module*

### **\$module purge**

*Will unload all modules*

### **\$module whatis <modulename>**

*Will show the module information*



# Examples

# SGE basics

```
#!/bin/bash  
# bigio_example.sh
```

```
# jobname  
#$ -N example
```

```
# use this shell  
#$ -S /bin/bash
```

```
# how much memory per processor  
#$ -l h_vmem=1g
```

```
# maximum run time  
#$ -l h_rt=00:05:00
```

```
# parallel environment SMP, 4 processes  
#$ -pe smp 4
```

```
# email user when done  
#$ -m me@myemail.com
```

```
runscript "$mytmp"
```

# “Big” I/O

- a. `mkdir mytemp @`
    - i. `/tmp`, or `/dev/shm`
  - b. Copy files into and untar
    - i. `cp remote/myfiles.tar mytempdir/`
    - ii. `tar xf myfiles.tar`
  - c. Run jobscript
    - i. with working output pointed to `mytempdir`
  - d. Tar contents of `mytempdir`
    - i. `tar cf myresults.tar myfiles`
  - e. Copy the results back to storage
    - i. `cp myresults.tar remote/myresults`
  - f. (check the move with checksum)
  - g. Cleanup (some systems do this for you)
    - i. `rm -r mytempdir`
-

# SGE

## Using the /tmp space

```
#!/bin/bash
#$ -N bigIO_example
#$ -S /bin/bash
#$ -l h_vmem=1g
#$ -l h_rt=02:00:00
#$ -pe smp 4
```startupscript
# mydir="$2" tmploc="$1"
starttar=$3 #
mytmp="$1"/"$2"

mkdir "$mytmp"
cp "$starttar" "$mytmp"
cd $mytmp
tar xf "$starttar"

runscript "$mytmp"

tar cf results.tar $mytmp
cp results.tar <remote system>
rm -r "$mytmp"

...`
```

# SGE

## Array jobs

```
#!/bin/bash  
#$ -N bigIO_example  
#$ -S /bin/bash  
#$ -l h_vmem=500m  
#$ -l h_rt=00:015:00  
#$ -pe smp 1  
#$ -t 1-5
```

e

# don't do this

# SGE

## Array jobs

1. Split the files into batched directories
2. Make a file listing those directories
  - a. `ls -l * > list.txt`
3. Get a count of files
  - a. `wc -l list.txt`
4. Use sed or awk to pop from list.txt and use bigIO pattern

```
#!/bin/bash
#$ -N bigIO_example
#$ -S /bin/bash
#$ -l h_vmem=500m
#$ -l h_rt=00:015:00
#$ -pe smp 1
#$ -t 1-<number from "wc -l">
```

```
infile=$(sed -n -e "$SGE_TASK_ID p" list.txt)
```

```
starttar="$infile"
mytmp="/tmp/"$2"
```

```
mkdir "$mytmp"
cp "$starttar" "$mytmp"
cd $mytmp
tar xf "$startar"
```

```
runscript "$mytmp"
```

```
tar cf "$infile"_results.tar $mytmp
cp "$infile"_results.tar <remote system>
rm -r "$mytemp"
```

# HPC resources

- With ITsupport at Einstein
- XSEDE
- Open Science Grid
- “The Cloud”
  - Amazon, Google, Microsoft and others